

Paper Summary Mip-NeRF


3. Method

A pixel's color is the integration of all incoming radiance within the pixel's frustum. NeRF casts a single infinitesimally narrow ray per pixel, resulting in aliasing. Mip-NeRF ameliorates this issue by casting a **cone** from each pixel.

Dividing the cone being cast into a series of **conical frustums**, and using **integrated** positional encoding (IPE)

3.1 Cone Tracing and Positional Encoding

Describing procedure in terms of an individual pixel of interest being rendered. Casting a cone from the camera's center of projection \mathbf{o} along the direction \mathbf{d} that passes through the pixel's center. The **radius** of the cone at the image plane $\mathbf{o} + \mathbf{d}$ is parameterized as \dot{r}



The set of positions \mathbf{x} that lie within a conical frustum between two t values $[t_0, t_1]$ is:

$$F(\mathbf{x}, \mathbf{o}, \mathbf{d}, \dot{r}, t_0, t_1) = 1 \left\{ \left(t_0 < \frac{\mathbf{d}^T(\mathbf{x} - \mathbf{o})}{\|\mathbf{d}\|_2^2} < t_1 \right) \wedge \left(\frac{\mathbf{d}^T(\mathbf{x} - \mathbf{o})}{\|\mathbf{d}\|_2 \|\mathbf{x} - \mathbf{o}\|_2} > \frac{1}{\sqrt{1 + (\dot{r}/\|\mathbf{d}\|_2)^2}} \right) \right\}$$

Prove is in the inference file.

Then we must construct a featurized representation of the volume inside this conical frustum. Simply compute the expected positional encoding of all coordinates that lie within the conical frustum:

$$\gamma^*(\mathbf{o}, \mathbf{d}, \dot{r}, t_0, t_1) = \frac{\int \gamma(\mathbf{x}) F(\mathbf{x}, \mathbf{o}, \mathbf{d}, \dot{r}, t_0, t_1) d\mathbf{x}}{\int F(\mathbf{x}, \mathbf{o}, \mathbf{d}, \dot{r}, t_0, t_1) d\mathbf{x}}$$

It's difficult on computation, therefore, approximate the **conical frustum** with a **multivariate Gaussian**, called **IPE**.

This Gaussian is fully characterized by three values (in addition to \mathbf{o} and \mathbf{d}): the mean distance along the ray μ_t , the variance along the ray σ_t^2 , and the variance perpendicular to the ray σ_r^2

.....

We can transfer this Gaussian from the coordinate frame of the conical frustum into world coordinates as follows:

$$\begin{aligned}\mu &= \mathbf{o} + \mu_t \mathbf{d} \\ \Sigma &= \sigma_t^2 (\mathbf{d} \mathbf{d}^T) + \sigma_r^2 \left(\mathbf{I} - \frac{\mathbf{d} \mathbf{d}^T}{\|\mathbf{d}\|_2^2} \right)\end{aligned}$$

give us final multivariate Gaussian, so here this is how we approximate the conical frustum.

Next, we derive the IPE, which is the expectation of a positionally-encoded coordinate distributed according to the previous Gaussian. Firstly, rewrite the PE in Equation 1 as Fourier Feature

$$\begin{aligned}& \begin{bmatrix} 1 & 0 & 0 & 2 & 0 & 0 & & 2^{L-1} & 0 & 0 \\ 0 & 1 & 0 & 0 & 2 & 0 & \dots & 0 & 2^{L-1} & 0 \\ 0 & 0 & 1 & 0 & 0 & 2 & & 0 & 0 & 2^{L-1} \end{bmatrix} \\ & \gamma(x) = \begin{bmatrix} \sin(\mathbf{P}x) \\ \cos(\mathbf{P}x) \end{bmatrix}\end{aligned}$$

This reparameterization allows us to derive a closed form for IPE. Using the fact $Cov[Ax, By] = ACov[x, y]B^T$, we can identify the mean and covariance of our conical frustum Gaussian

$$\begin{aligned}\mu_\gamma &= \mathbf{P}\mu \\ \Sigma_\gamma &= \mathbf{P}\Sigma\mathbf{P}^T\end{aligned}$$

Final Step: computing the expectations over this lifted multivariate Gaussian, modulated by the sine and cosine of position. These expectations have simple closed-form expressions:

$$\begin{aligned}E_{x \sim N(\mu, \sigma^2)}[\sin(x)] &= \sin(\mu) \exp\left(-\left(\frac{1}{2}\right)\sigma^2\right) \\ E_{x \sim N(\mu, \sigma^2)}[\cos(x)] &= \cos(\mu) \exp\left(-\left(\frac{1}{2}\right)\sigma^2\right)\end{aligned}$$

With this we can compute our final IPE features as the expected sines and cosines of the mean and the diagonal of the covariance matrix:

$$\begin{aligned}\gamma(\mu, \Sigma) &= E_{x \sim N(\mu_\gamma, \Sigma_\gamma)}[\gamma(x)] \\ &= \begin{bmatrix} \sin(\mu_\gamma) \circ \exp\left(-\left(\frac{1}{2}\right) \text{diag}(\Sigma_\gamma)\right) \\ \cos(\mu_\gamma) \circ \exp\left(-\left(\frac{1}{2}\right) \text{diag}(\Sigma_\gamma)\right) \end{bmatrix}\end{aligned}$$